

# 第17组 云原生大作业说明文档

云原生斗地主 GROUP nju17

项目地址:

 [github](#)

 [gitee](#)

## 小组成员

姓名	学号
张铭铭	211250234
胡家睿	211250020
宋毅恒	211250022

## 功能要求

### 实现接口和限流功能

在项目中新建一个Controller，实现Rest接口如下：

```
@RestController
public class DemoController {

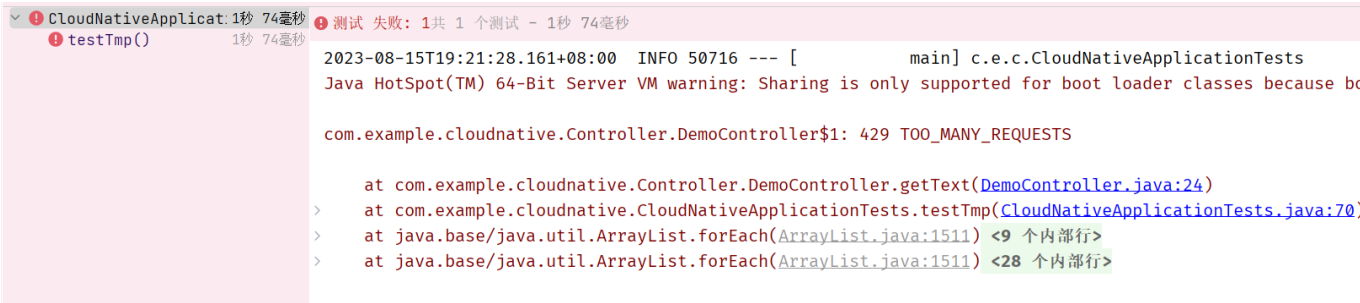
    private final RateLimiter rateLimiter = RateLimiter.create(100.0);

    @GetMapping("/api/text")
    @ResponseStatus(HttpStatus.OK)
    @ResponseBody
    public String getText(){
        if(!rateLimiter.tryAcquire(1))
            throw new HttpStatusCodeException(HttpStatus.TOO_MANY_REQUESTS) {
        };
        return "{\"name\":\"云原生斗地主\",\"number\":\"nju17\"}";
    }

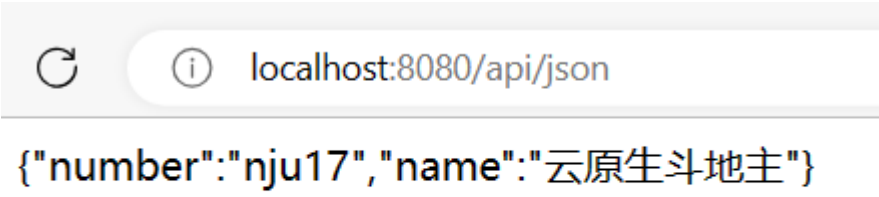
    @GetMapping( "/api/json")
    @ResponseStatus(HttpStatus.OK)
    @ResponseBody
    public String getJson() {
        if(!rateLimiter.tryAcquire(1)) {
            throw new HttpStatusCodeException(HttpStatus.TOO_MANY_REQUESTS) {};
        }
        JSONObject json = new JSONObject();
        try {
            json.put("name", "云原生斗地主");
        }
    }
}
```

```
        json.put("number", "nju17");
    } catch (JSONException e) {
        throw new RuntimeException(e);
    }
    String ret = json.toString();
    return ret;
}
}
```

其中，限流功能使用了RateLimiter相关接口来实现，限制每秒最多处理100个请求。如果请求过于频繁，则会返回429错误，如下所示（使用springboot test测试）：



本地运行后，该接口可以通过访问<http://localhost:8080/api/json> 或 <http://localhost:8080/api/text> 来测试可用性：



### 实现Prometheus监控

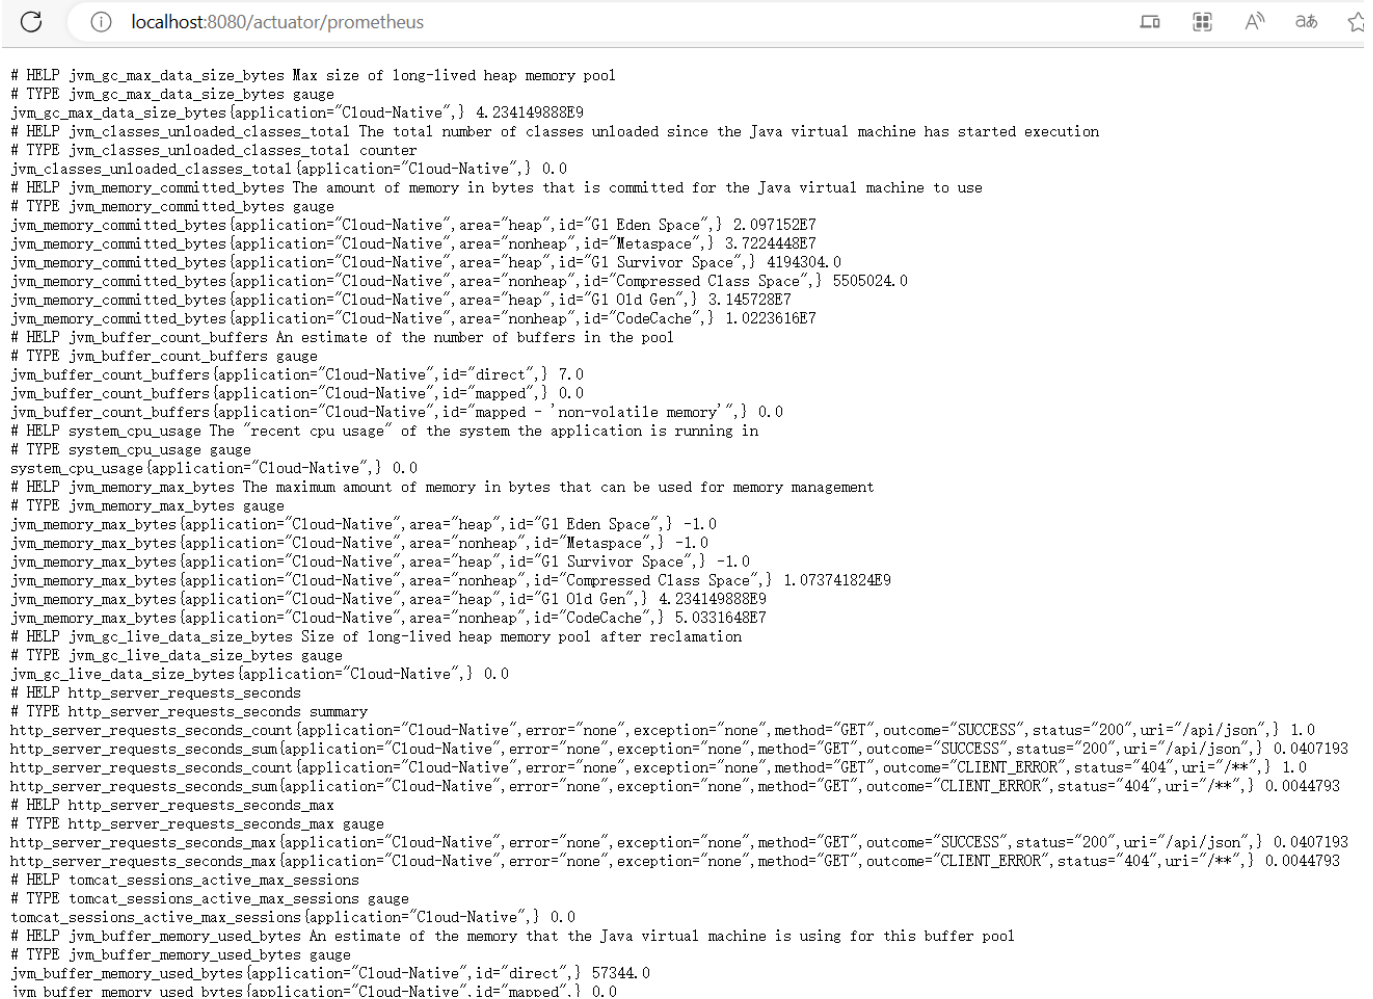
在项目的application.properties中添加如下配置：



并在pom.xml中添加相关依赖：

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-registry-prometheus</artifactId>
</dependency>
```

本地运行后，即可在<http://localhost:8080/actuator/prometheus> 中查看到相关监控信息：



```
# HELP jvm_gc_max_data_size_bytes Max size of long-lived heap memory pool
# TYPE jvm_gc_max_data_size_bytes gauge
jvm_gc_max_data_size_bytes{application="Cloud-Native",} 4.234149888E9
# HELP jvm_classes_unloaded_classes_total The total number of classes unloaded since the Java virtual machine has started execution
# TYPE jvm_classes_unloaded_classes_total counter
jvm_classes_unloaded_classes_total{application="Cloud-Native",} 0.0
# HELP jvm_memory_committed_bytes The amount of memory in bytes that is committed for the Java virtual machine to use
# TYPE jvm_memory_committed_bytes gauge
jvm_memory_committed_bytes{application="Cloud-Native",area="heap",id="G1 Eden Space",} 2.097152E7
jvm_memory_committed_bytes{application="Cloud-Native",area="nonheap",id="Metaspace",} 3.7224448E7
jvm_memory_committed_bytes{application="Cloud-Native",area="heap",id="G1 Survivor Space",} 4194304.0
jvm_memory_committed_bytes{application="Cloud-Native",area="nonheap",id="Compressed Class Space",} 5505024.0
jvm_memory_committed_bytes{application="Cloud-Native",area="heap",id="G1 Old Gen",} 3.145728E7
jvm_memory_committed_bytes{application="Cloud-Native",area="nonheap",id="CodeCache",} 1.0223616E7
# HELP jvm_buffer_count_buffers An estimate of the number of buffers in the pool
# TYPE jvm_buffer_count_buffers gauge
jvm_buffer_count_buffers{application="Cloud-Native",id="direct",} 7.0
jvm_buffer_count_buffers{application="Cloud-Native",id="mapped",} 0.0
jvm_buffer_count_buffers{application="Cloud-Native",id="mapped - 'non-volatile memory'",} 0.0
# HELP system_cpu_usage The "recent cpu usage" of the system the application is running in
# TYPE system_cpu_usage gauge
system_cpu_usage{application="Cloud-Native",} 0.0
# HELP jvm_memory_max_bytes The maximum amount of memory in bytes that can be used for memory management
# TYPE jvm_memory_max_bytes gauge
jvm_memory_max_bytes{application="Cloud-Native",area="heap",id="G1 Eden Space",} -1.0
jvm_memory_max_bytes{application="Cloud-Native",area="nonheap",id="Metaspace",} -1.0
jvm_memory_max_bytes{application="Cloud-Native",area="heap",id="G1 Survivor Space",} -1.0
jvm_memory_max_bytes{application="Cloud-Native",area="nonheap",id="Compressed Class Space",} 1.073741824E9
jvm_memory_max_bytes{application="Cloud-Native",area="heap",id="G1 Old Gen",} 4.234149888E9
jvm_memory_max_bytes{application="Cloud-Native",area="nonheap",id="CodeCache",} 5.0331648E7
# HELP jvm_gc_live_data_size_bytes Size of long-lived heap memory pool after reclamation
# TYPE jvm_gc_live_data_size_bytes gauge
jvm_gc_live_data_size_bytes{application="Cloud-Native",} 0.0
# HELP http_server_requests_seconds
# TYPE http_server_requests_seconds summary
http_server_requests_seconds_count{application="Cloud-Native",error="none",exception="none",method="GET",outcome="SUCCESS",status="200",uri="/api/json",} 1.0
http_server_requests_seconds_sum{application="Cloud-Native",error="none",exception="none",method="GET",outcome="SUCCESS",status="200",uri="/api/json",} 0.0407193
http_server_requests_seconds_count{application="Cloud-Native",error="none",exception="none",method="GET",outcome="CLIENT_ERROR",status="404",uri="/**",} 1.0
http_server_requests_seconds_sum{application="Cloud-Native",error="none",exception="none",method="GET",outcome="CLIENT_ERROR",status="404",uri="/**",} 0.0044793
# HELP http_server_requests_seconds_max
# TYPE http_server_requests_seconds_max gauge
http_server_requests_seconds_max{application="Cloud-Native",error="none",exception="none",method="GET",outcome="SUCCESS",status="200",uri="/api/json",} 0.0407193
http_server_requests_seconds_max{application="Cloud-Native",error="none",exception="none",method="GET",outcome="CLIENT_ERROR",status="404",uri="/**",} 0.0044793
# HELP tomcat_sessions_active_max_sessions
# TYPE tomcat_sessions_active_max_sessions gauge
tomcat_sessions_active_max_sessions{application="Cloud-Native",} 0.0
# HELP jvm_buffer_memory_used_bytes An estimate of the memory that the Java virtual machine is using for this buffer pool
# TYPE jvm_buffer_memory_used_bytes gauge
jvm_buffer_memory_used_bytes{application="Cloud-Native",id="direct",} 57344.0
jvm_buffer_memory_used_bytes{application="Cloud-Native",id="mapped",} 0.0
```

## 统一限流

统一限流暂未实现~

## DevOps 要求

## Dockerfile与K8s容器编排

Dockerfile

```
Dockerfile x
1 FROM openjdk:17
2
3 LABEL authors="nju17"
4
5 COPY ./release/Cloud-Native-0.0.1-SNAPSHOT.jar /app/Cloud-Native.jar
6
7 WORKDIR /app
8
9 EXPOSE 8080
10
11 ENTRYPOINT ["java", "-jar", "Cloud-Native.jar"]
```

deployment.yaml

```
1 apiVersion: apps/v1 #api版本
2 kind: Deployment
3 metadata:
4   labels: #资源标签
5     app: cloud-native
6   name: cloud-native #资源名
7   namespace: nju17 #资源部署的名空间
8 spec:
9   replicas: 3 #副本数目
10  strategy:
11    type: RollingUpdate #滚动更新策略
12    rollingUpdate:
13      maxSurge: 25% #最大额外副本
14      maxUnavailable: 25% #更新中进入不可用状态Pod的最大值
15  selector:
16    matchLabels:
17      app: cloud-native #匹配标签
18  template:
19    metadata:
20      annotations: #自定义注解，接入Prometheus
21        prometheus.io/path: /actuator/prometheus
22        prometheus.io/port: "8080"
23        prometheus.io/scheme: http
24        prometheus.io/scrape: "true"
25      labels:
26        app: cloud-native #资源标签
27    spec:
28      containers:
29        - image: harbor.edu.cn/nju17/cloud-native:9 #镜像地址
30          name: cloud-native #容器名
31  ---
32  apiVersion: v1 #api版本
33  kind: Service
34  metadata:
35    name: cloud-native #资源名
36  labels:
37    app: cloud-native #资源标签
38  spec:
39    type: NodePort
40    selector:
41      app: cloud-native
42    ports:
43      - name: tcp8080 #端口名
44        protocol: TCP #TCP协议
45        port: 8080 #service端口
46        targetPort: 8080 #容器暴露的端口
```

相关注解在代码注释中。截图中的代码适配下述方案二流水线，在Jenkins部分会进行说明。

下述部分是另外一套文件，适配方案一流水线：

- Dockerfile

```

1  FROM openjdk:17
2
3  LABEL authors="nju17"
4
5  COPY ./target/Cloud-Native-0.0.1-SNAPSHOT.jar /app/Cloud-Native.jar
6
7  WORKDIR /app
8
9  EXPOSE 8080
10
11 ENTRYPOINT ["java", "-jar", "Cloud-Native.jar"]

```

- deployment.yaml 仅修改图示部分

```

labels:
  app: cloud-native #资源标签
spec:
  containers:
  - image: harbor.edu.cn/nju17/cloud-native:{VERSION} #镜像地址
    name: cloud-native #容器名
---
apiVersion: v1 #api版本
kind: Service
metadata:
  name: cloud-native #资源名

```

## Jenkins

由于软院Jenkins服务器的master节点一直被占用，因此我们原本做好的Jenkin流水线无法运行，后来临时做了另一条流水线。结果后面master节点又正常工作，又成功运行了原本的方案。两套方案如下：

- 方案一：原本的Jenkin流水线方案，同时使用了master节点和slave节点进行构建，其中master节点完成了持续集成的功能，slave节点完成了持续部署的功能。流水线如下：

```

pipeline {
  agent none
  stages {
    stage('Clone Code') {
      agent {
        label 'master'
      }
      steps {
        echo "1.Clone From Gitee"
        sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
          'username=' + '211250234' +
          '&' +
          'password=' + 'xxxxxxx' + '"' //密码已隐藏

```

```

        git url: 'https://gitee.com/irisalt/cloud-native.git', branch:
'main'
    }
}

stage('Maven Build') {
    agent {
        docker {
            image 'maven:latest'
            args ' -v /root/.m2:/root/.m2'
        }
    }
    steps {
        echo "2. Using Maven to Build"
        sh 'mvn -B clean package'
    }
}

stage('Build Image') {
    agent {
        label 'master'
    }
    steps {
        echo "3. Build Image"
        sh 'docker build -t cloud-native:${BUILD_ID} .'
        sh 'docker tag cloud-native:${BUILD_ID} harbor.edu.cn/nju17/cloud-
native:${BUILD_ID}'
    }
}

stage('Push Image') {
    agent {
        label 'master'
    }
    steps {
        echo "4. Push Docker Image"
        sh 'docker login harbor.edu.cn ' +
            '-u ' + 'nju17' +
            ' -p ' + 'nju172023'
        sh 'docker push harbor.edu.cn/nju17/cloud-native:${BUILD_ID}'
    }
}
}

node('slave') {
    container('jnlp-kubect1') {
        stage('Clone & Change YAML') {
            echo "5. Clone YAML to Slave and Change YAML"
            //xxx needs to be replaced
            sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
                'username=' + '211250234' +
                '&' +
                'password=' + 'xxxxxxx' + '"' //密码已隐藏

```

```

        git url: 'https://gitee.com/irisalt/cloud-native.git', branch: 'main'
        sh 'sed -i "s#{VERSION}#{BUILD_ID}#g" deployment.yaml'
    }

    stage ('Deploy') {
        echo "6. Deploy to K8s"
        //xxx needs to be replaced
        sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
            'username=' + '211250234' +
            '&' +
            'password=' + 'xxxxxxx' + '"' //密码已隐藏
        sh 'docker login harbor.edu.cn ' +
            '-u ' + 'nju17' +
            '-p ' + 'nju172023'
        sh 'docker pull harbor.edu.cn/nju17/cloud-native:${BUILD_ID}'
        sh 'kubectl apply -f deployment.yaml -n nju17'
    }

    stage('Monitor') {
        echo "7. Start Monitor"
        sh 'kubectl apply -f monitor.yaml -n monitoring'
    }
}
}

```

slave节点完成了从镜像仓库拉取镜像，部署到K8s集群的任务。在部署时，会自动修改deployment.yaml文件中的镜像TAG，以实现持续部署的功能。

master节点完成了代码上传，镜像构建，上传到镜像仓库的任务。此外在Maven构建时，已经通过了本地写好的单元测试，如下所示：

```

java hotspot(TM) 64-bit server vm warning. Shading is only supported for boot loader classes bec
[INFO] Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 5.6 s - in com.example.cl
[INFO]
[INFO] Results:
[INFO]
[INFO] Tests run: 4, Failures: 0, Errors: 0, Skipped: 0
[INFO]
[INFO]
[INFO] --- jar:3.3.0:jar (default-jar) @ Cloud-Native ---
[INFO] Building jar: F:\Cloud-Native\target\Cloud-Native-0.0.1-SNAPSHOT.jar
[INFO]
[INFO] --- spring-boot:3.1.2:repackage (repackage) @ Cloud-Native ---
[INFO] Replacing main artifact F:\Cloud-Native\target\Cloud-Native-0.0.1-SNAPSHOT.jar with repac
[INFO] The original artifact has been renamed to F:\Cloud-Native\target\Cloud-Native-0.0.1-SNAPS
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 13.129 s

```

编写的单元测试代码如下：

```

@SpringBootTest
class CloudNativeApplicationTests {

```

```
private DemoController demoController = new DemoController();
private static final String EXPECTED_TEXT = "{\"name\":\"云原生斗地主\"," +
    "\",\"number\":\"nju17\"}";
@BeforeEach
void initAll() {demoController = new DemoController();}
@Test
void testGetText() { // 测试可用性
    String result = demoController.getText();
    assertEquals(EXPECTED_TEXT);
}
@Test
void test429(){ // 测试限流功能, 应该爆429并且通过测试
    try {
        for(int i = 0; i < 100; i++) {
            Thread.sleep(5);
            demoController.getText();
        }
        assertFalse();
    } catch (Exception e) {
        assertEquals(e.getMessage(), "429 TOO_MANY_REQUESTS");
    }
}
@Test
void test429Two(){ // 测试限流功能, 应该爆429并且通过测试
    try {
        for(int i = 0; i < 100; i++) {
            demoController.getText();
        }
        assertFalse();
    } catch (Exception e) {
        assertEquals(e.getMessage(), "429 TOO_MANY_REQUESTS");
    }
}
@Test
void testEdge(){ // 边界测试, 应该表现为不会爆429
    try {
        for(int i = 0; i < 100; i++) {
            Thread.sleep(11);
            demoController.getText();
        }
        assertTrue();
    } catch (Exception e) {
        if(e.getMessage().equals("429 TOO_MANY_REQUESTS")){
            assertFalse();
        };
        assertTrue();
    }
}
```



流水线构建结果如下：

Dashboard017Cloud-Native

立即构建

配置

删除 Pipeline

完整阶段视图

GitHub

打开 Blue Ocean

重命名

流水线语法

GitHub Hook Log

Build History构建历史

find

#142023-8-15 下午2:20

#132023-8-15 下午2:20

#122023-8-15 下午2:20

#112023-8-15 下午2:06

#102023-8-15 下午2:02

阶段视图

Average stage times:  
(Average full run time: ~57s)

	Clone Code	Maven Build	Build Image	Push Image	Clone & Change YAML	Deploy	Monitor
#14	1s	14s	1s	936ms	16s	4s	1s
#13	3s	42ms	46ms	47ms			
#12	1s	8s	40ms	41ms			
#11	1s	13s	1s	937ms	15s	4s	969ms
#10	1s	14s	1s	2s	17s	5s	1s

• [nju17@host-172-29-4-18 ~]\$ kubectl get pods -n nju17

NAMEREADYSTATUSRESTARTSAGE

cloud-native-866946c48c-d74t41/1Running036s

cloud-native-866946c48c-nl2cb1/1Running036s

cloud-native-866946c48c-wttzw1/1Running036s

• [nju17@host-172-29-4-18 ~]\$ kubectl get service -n nju17

NAMECLUSTER-IPEXTERNAL-IPPORT(S)AGE

cloud-native10.103.149.215<none>8080:30713/TCP39s

○ [nju17@host-172-29-4-18 ~]\$

- 方案二：由于准备开始构建流水线的时候master节点都用不了，因此得修改流水线。这是因为master节点和slave节点的功能是不一样的，经过本人测试，在slave节点上我们无法使用mvn docker等命令（应该 是没有预装相关的环境）。因此小组同时使用了另一套方案：
  - 在本地手动将docker镜像上传到镜像仓库，并手动指定TAG：

```
ttha@LAPTOP-IRV0R5A3:~/Cloud-Native$ docker login harbor.edu.cn
Username: nju17
Password:
WARNING! Your password will be stored unencrypted in /home/ttha/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
ttha@LAPTOP-IRV0R5A3:~/Cloud-Native$ docker build cloud-native:9 .
ERROR: "docker buildx build" requires exactly 1 argument.
See 'docker buildx build --help'.

Usage:  docker buildx build [OPTIONS] PATH | URL | -

Start a build
ttha@LAPTOP-IRV0R5A3:~/Cloud-Native$ docker build -t cloud-native:9 .
[+] Building 31.6s (8/8) FINISHED
=> [internal] load build definition from Dockerfile                                0.0s
=> => transferring dockerfile: 235B                                              0.0s
=> [internal] load .dockerignore                                                  0.0s
=> => transferring context: 2B                                                    0.0s
=> [internal] load metadata for docker.io/library/openjdk:17                    31.4s
=> [internal] load build context                                                  0.0s
=> => transferring context: 90B                                                  0.0s
=> [1/3] FROM docker.io/library/openjdk:17@sha256:528707081fdb9562eb819128a9f85ae7fe000e2fbaeaf9f87662e7b3f38cb7  0.0s
=> CACHED [2/3] COPY ./release/Cloud-Native-0.0.1-SNAPSHOT.jar /app/Cloud-Native.jar 0.0s
=> CACHED [3/3] WORKDIR /app                                                      0.0s
=> exporting to image                                                            0.0s
=> => exporting layers                                                            0.0s
=> => writing image sha256:c58d5bc28e3708feeced81cb5e73fa54680f2c373ff4aaf861fc8c717915f0b9 0.0s
=> => naming to docker.io/library/cloud-native:9                                0.0s
ttha@LAPTOP-IRV0R5A3:~/Cloud-Native$ docker tag cloud-native:9 harbor.edu.cn/nju17/cloud-native:9
ttha@LAPTOP-IRV0R5A3:~/Cloud-Native$ docker push harbor.edu.cn/nju17/cloud-native:9
The push refers to repository [harbor.edu.cn/nju17/cloud-native]
5f70bf18a086: Pushed
aaef6f773b7c: Pushed
dc9fa3d8b576: Pushed
27ee19dc88f2: Pushed
c8dd97366670: Pushed
9: digest: sha256:63e943e866a49e420dc1becaa08d659cdb866f83fcc386c9852cbe5be881608b size: 1372
```

9 / 19

- 构建流水线如下，该流水线实际上实现了持续部署的功能：

```

pipeline {
  agent none
  stages {
    stage('Clone Code') {
      agent {
        label 'slave'
      }
      steps {
        echo "1.Clone From Gitee"
        //xxx needs to be replaced
        sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
            'username=' + '211250234' +
            '&' +
            'password=' + 'xxxxxxx' + '"' //密码已隐藏
        git url: 'https://gitee.com/irisalt/cloud-native.git',
        branch: 'main'
      }
    }
  }
}

node('slave') {
  container('jnlp-kubect1') {
    stage('Clone & Change YAML') {
      echo "2. Clone YAML to Slave and Change YAML"
      //xxx needs to be replaced
      sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
          'username=' + '211250234' +
          '&' +
          'password=' + 'xxxxxxx' + '"' //密码已隐藏
      git url: 'https://gitee.com/irisalt/cloud-native.git', branch: 'main'
    }

    stage('Deploy') {
      echo "3. Deploy to K8s"
      //xxx needs to be replaced
      sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
          'username=' + '211250234' +
          '&' +
          'password=' + 'xxxxxxx' + '"' //密码已隐藏
      sh 'docker login harbor.edu.cn ' +
          '-u ' + 'nju17' +
          ' -p ' + 'nju172023'
      sh 'docker pull harbor.edu.cn/nju17/cloud-native:9'
      // sh 'kubectl delete deployment cloud-native -n nju17'
      sh 'kubectl apply -f deployment.yaml -n nju17'
      // sh 'kubectl scale deployment cloud-native --replicas 1 -n
nju17'
    }

    stage('Monitor') {

```

```
    echo "4. Start Monitor"
    sh 'kubectl apply -f monitor.yaml -n monitoring'
  }
}
```

实际运行的情况如下，访问的url为<http://172.29.4.18:32510/>：

017Cloud-Native-backup ☆ ⚙

活动 分支 Pull Requests

运行

Disable

状态	运行	提交	消息	持续时间	完成
✔	32	—	由用户 nju17 启动	1m 7s	3 hours ago ↺
✖	31	—	由用户 nju17 启动	1m 0s	3 hours ago ↺
✖	30	—	由用户 nju17 启动	59s	3 hours ago ↺
✔	29	—	Manual capacity expansion 1→3	59s	4 hours ago ↺
✔	28	—	基于#26 回放	59s	9 hours ago ↺
✔	26	—	由用户 nju17 启动	59s	16 hours ago ↺
✔	25	—	由用户 nju17 启动	1m 6s	16 hours ago ↺
✖	24	—	have a try again	1m 4s	16 hours ago ↺
✔	23	—	由用户 nju17 启动	53s	16 hours ago ↺

Tolerations:

node.kubernetes.io/not-ready:NoExecute op=Exists for 300s  
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:

<none>

• [nju17@host-172-29-4-18 ~]\$ kubectl get pods -n nju17

NAME READY STATUS RESTARTS AGE

cloud-native-866946c48c-55bjv 1/1 Running 0 3m46s

• [nju17@host-172-29-4-18 ~]\$ kubectl get pods -n nju17

NAME READY STATUS RESTARTS AGE

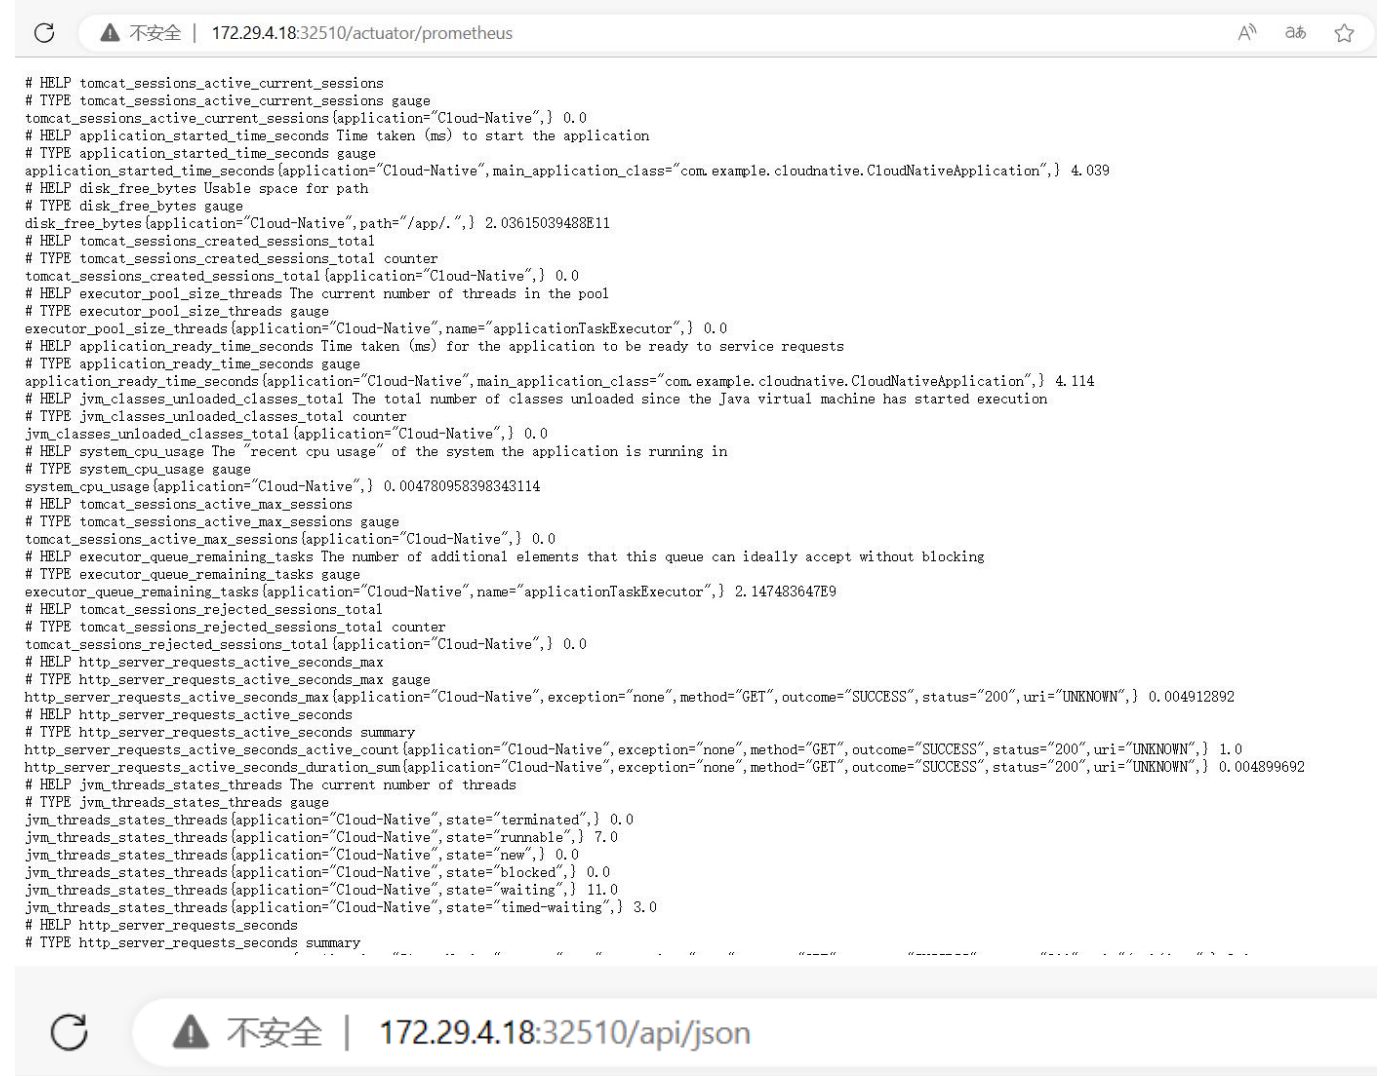
cloud-native-866946c48c-55bjv 1/1 Running 0 4m45s

• [nju17@host-172-29-4-18 ~]\$ kubectl get service -n nju17

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE

cloud-native NodePort 10.107.125.168 <none> 8080:32510/TCP 26s

○ [nju17@host-172-29-4-18 ~]\$



扩容场景

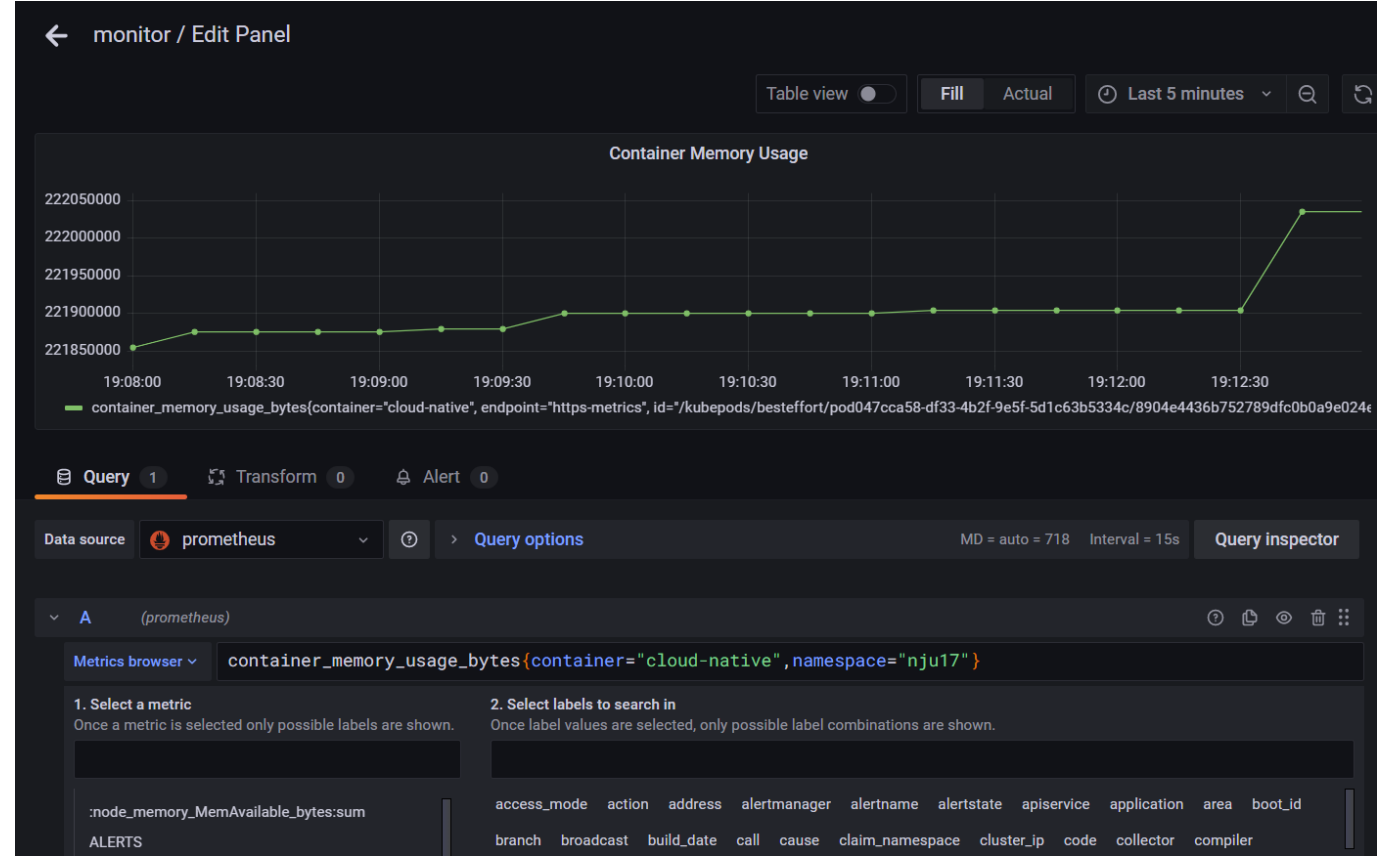
Prometheus metrics接口

配置一个ServiceMonitor，用于监控应用的metrics接口，访问的url为  
<http://172.29.4.18:32510/actuator/prometheus>，配置如下图所示：

```
1  apiVersion: monitoring.coreos.com/v1
2  kind: ServiceMonitor
3  metadata:
4    labels:
5      k8s-app: cloud-native
6      name: cloud-native
7      namespace: monitoring
8  spec:
9    endpoints:
10     - interval: 30s #监控数据抓取间隔设为30s
11       port: tcp8080 #端口名
12       path: /actuator/prometheus
13       scheme: 'http' #接口协议
14    selector:
15      matchLabels:
16        app: cloud-native #监控目标Service的标签
17      namespaceSelector:
18        matchNames:
19        - nju17
```

Grafana监控

流水线部署完成后，在软件研发效能支撑平台的grafana平台上可以直接查询到对应的容器和名空间。通过可视化工具可以直接创建所需的图表（如下图）。



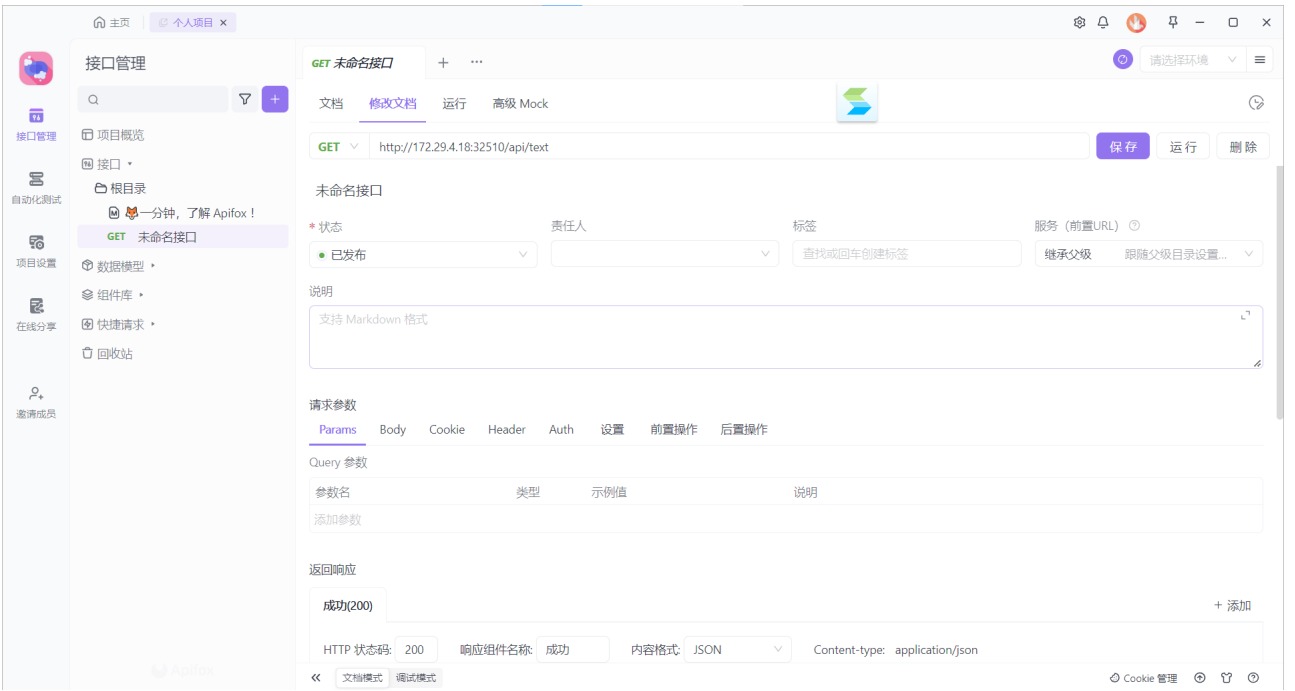
通过此方式创建CPU、内存、JVM的空间使用图表（如下图）。



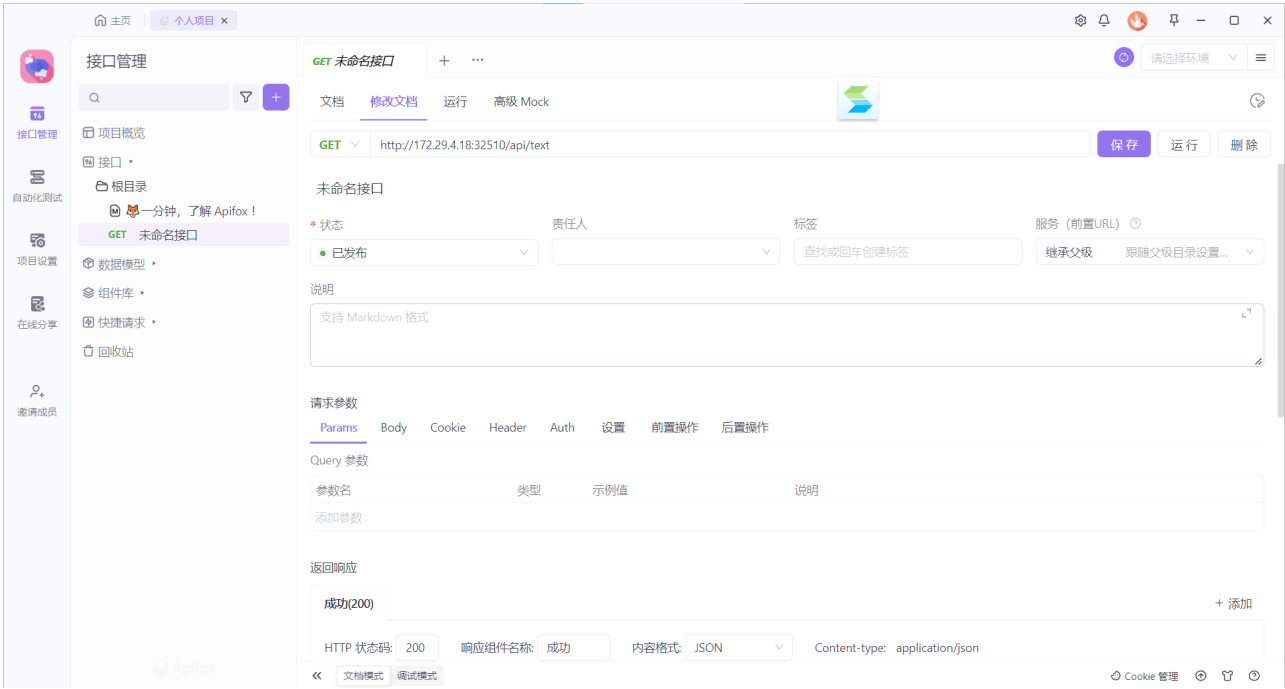
压力测试

- 使用Apifox进行压力测试

在Apifox中添加实现的接口（如下图）。

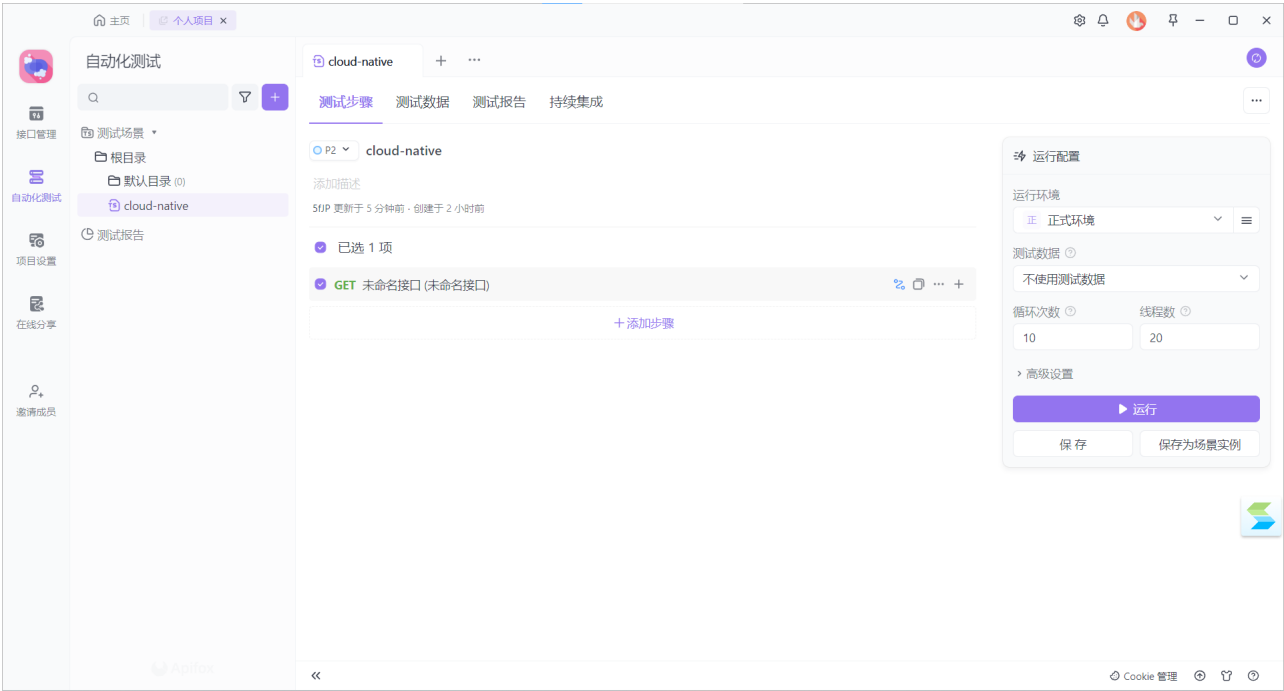


使用Apifox的自动化测试功能,设置循环次数为10次，线程数为20(如下图)。



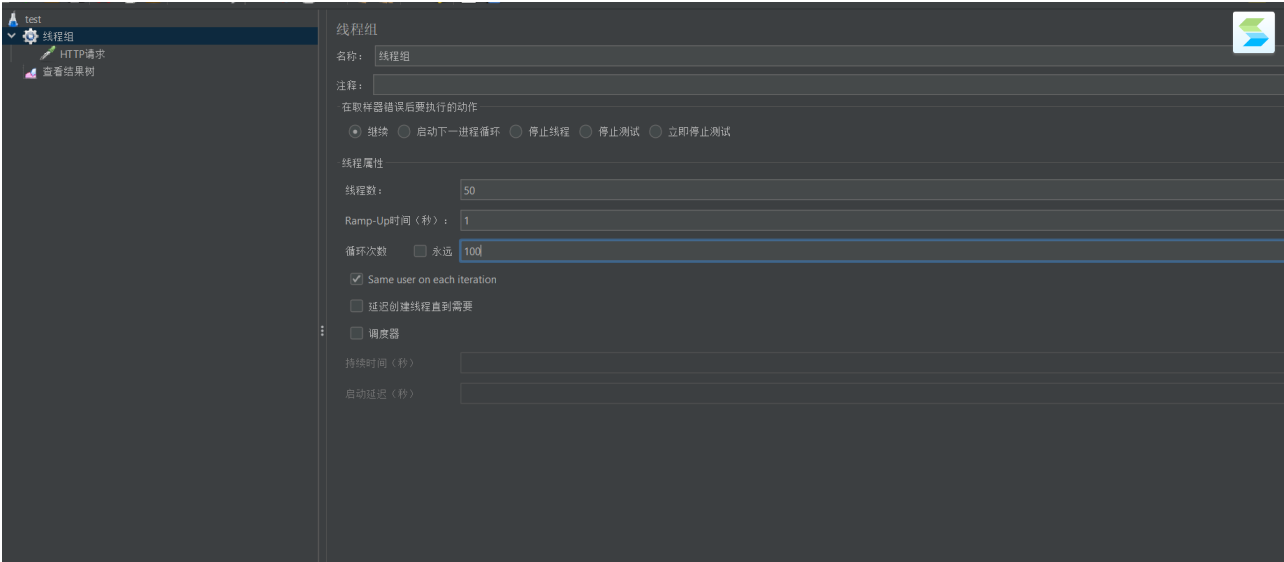
测试完成后可以查看grafana中容器内存使用有明显上升（由于网络限制测试时间花费较长，故没有触发限流）。





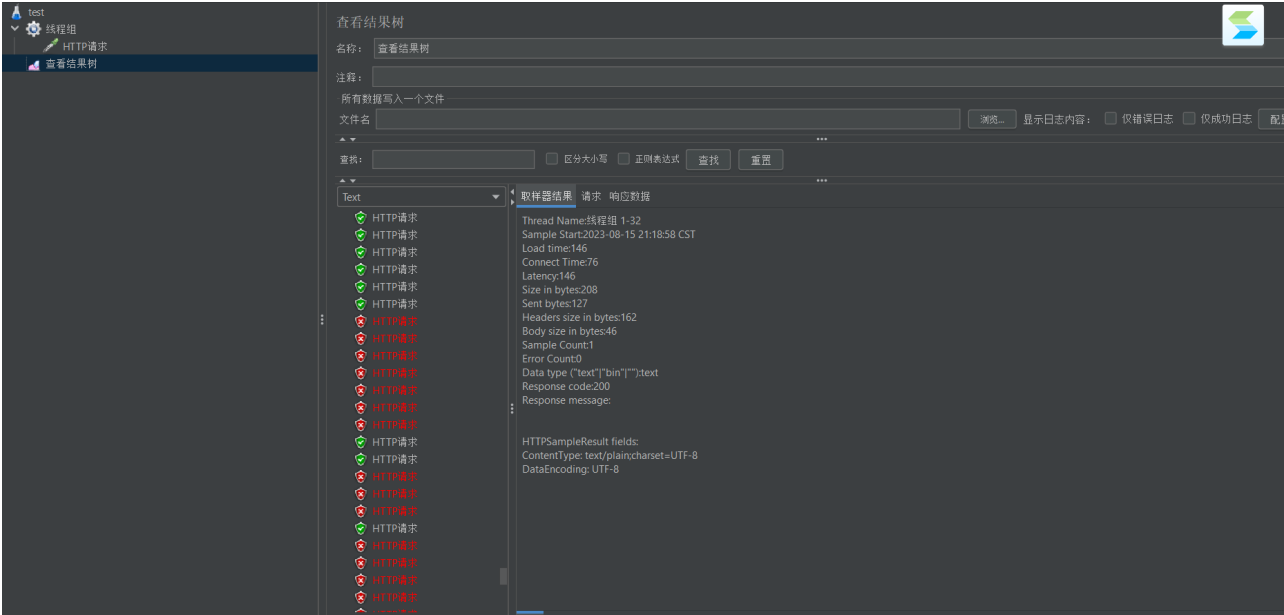
- 使用Apache JMeter进行压力测试

由于ApiFox的测试时间较长，没有达到理想的效果，故再采用Apache JMeter进行压力测试。

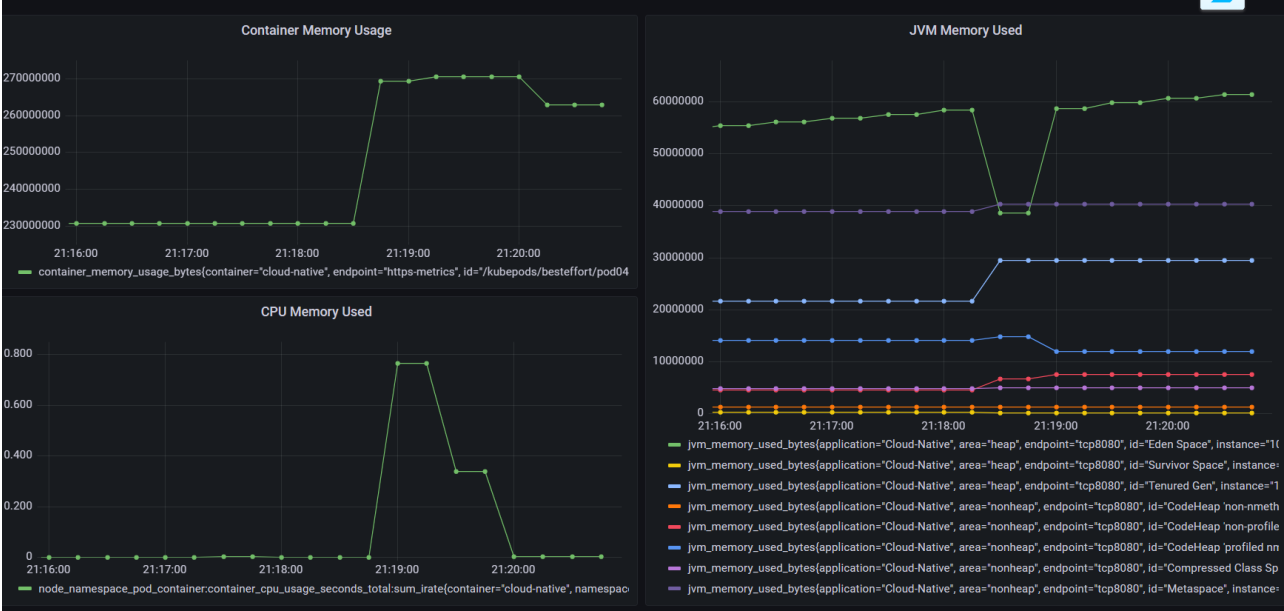




测试完成后可以看到出现了429错误，说明限流功能正常。



查看监控大屏，大屏出现了明显的资源占用上升：



手工扩容

可以修改deployment.yaml文件中replicas的值，再次部署。此处采取修改流水线的方式，直接修改Deploy阶段，增加对replicas的修改：

```
// 其余部分省流

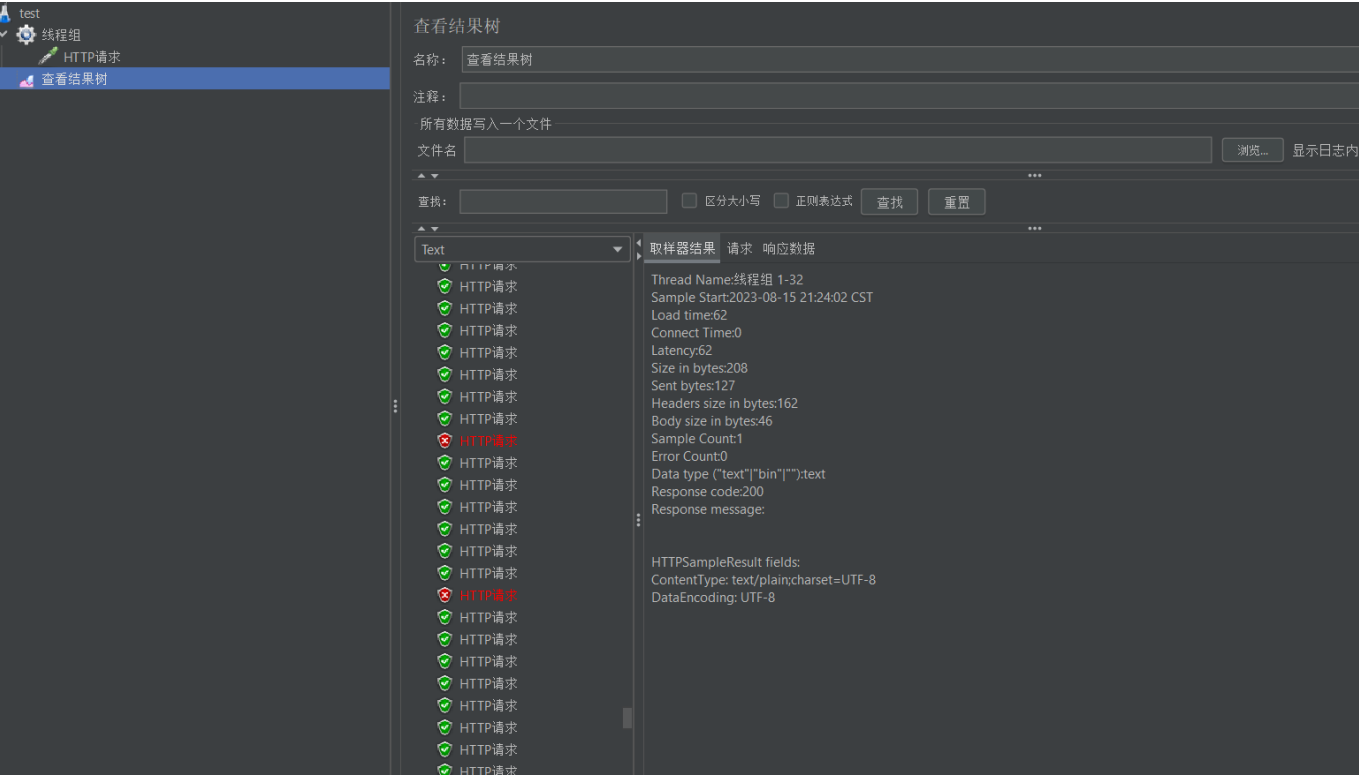
stage ('Deploy') {
    echo "3. Deploy to K8s"
    sh 'curl "http://p.nju.edu.cn/portal_io/login?" +
        'username=' + '211250234' +
        '&' +
        'password=' + 'xxxxxxx' + '"' //密码已隐藏
    sh 'docker login harbor.edu.cn ' +
        '-u ' + 'nju17' +
        '-p ' + 'nju172023'
```

```
sh 'docker pull harbor.edu.cn/nju17/cloud-native:9'
// sh 'kubectl delete deployment cloud-native -n nju17'
sh 'kubectl apply -f deployment.yaml -n nju17'
sh 'kubectl scale deployment cloud-native --replicas 3 -n nju17'
}
```

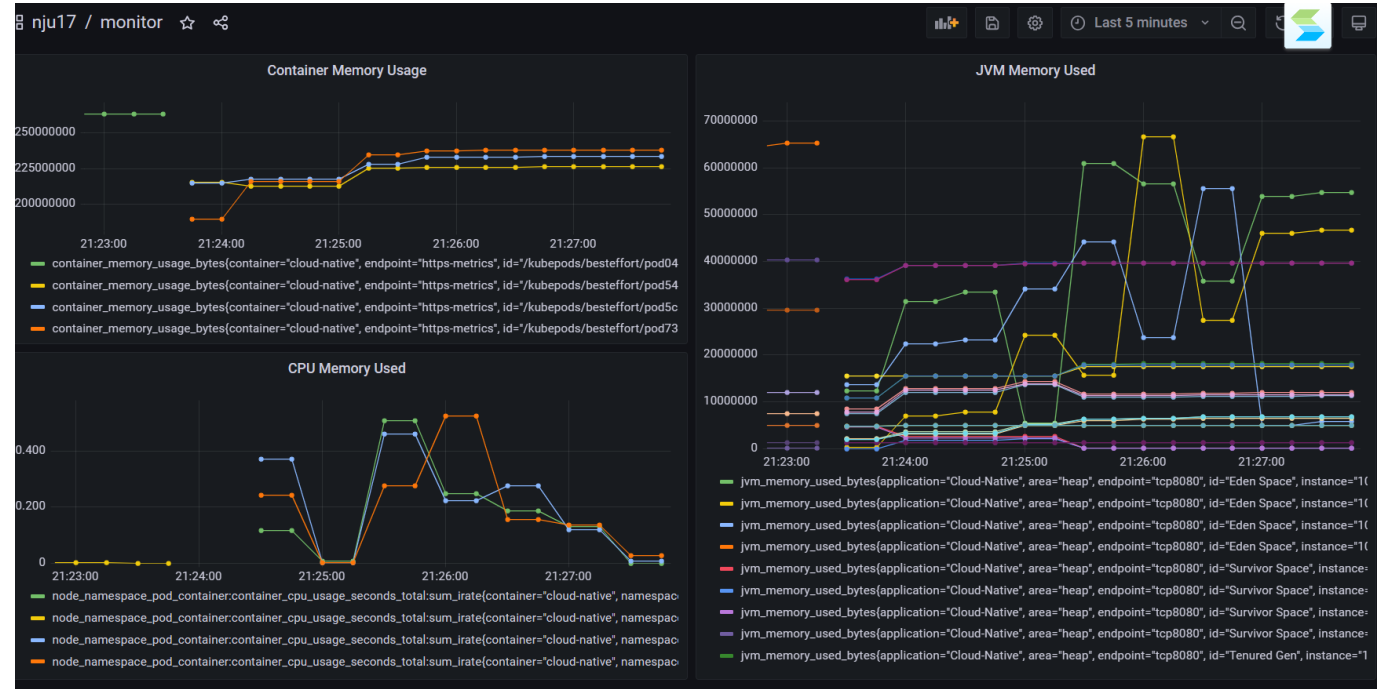
增加sh 'kubectl scale deployment cloud-native --replicas 3 -n nju17'即可在部署时实现扩容，扩容结果如下。

```
[nju17@host-172-29-4-18 ~]$ kubectl get pods -n nju17
NAME                                READY   STATUS    RESTARTS   AGE
cloud-native-866946c48c-jqm8l      1/1     Running   0          3h26m
[nju17@host-172-29-4-18 ~]$ kubectl get pods -n nju17
NAME                                READY   STATUS    RESTARTS   AGE
cloud-native-866946c48c-c7559      1/1     Running   0          12s
cloud-native-866946c48c-p2xxl      1/1     Running   0          12s
cloud-native-866946c48c-v55bt      1/1     Running   0          12s
```

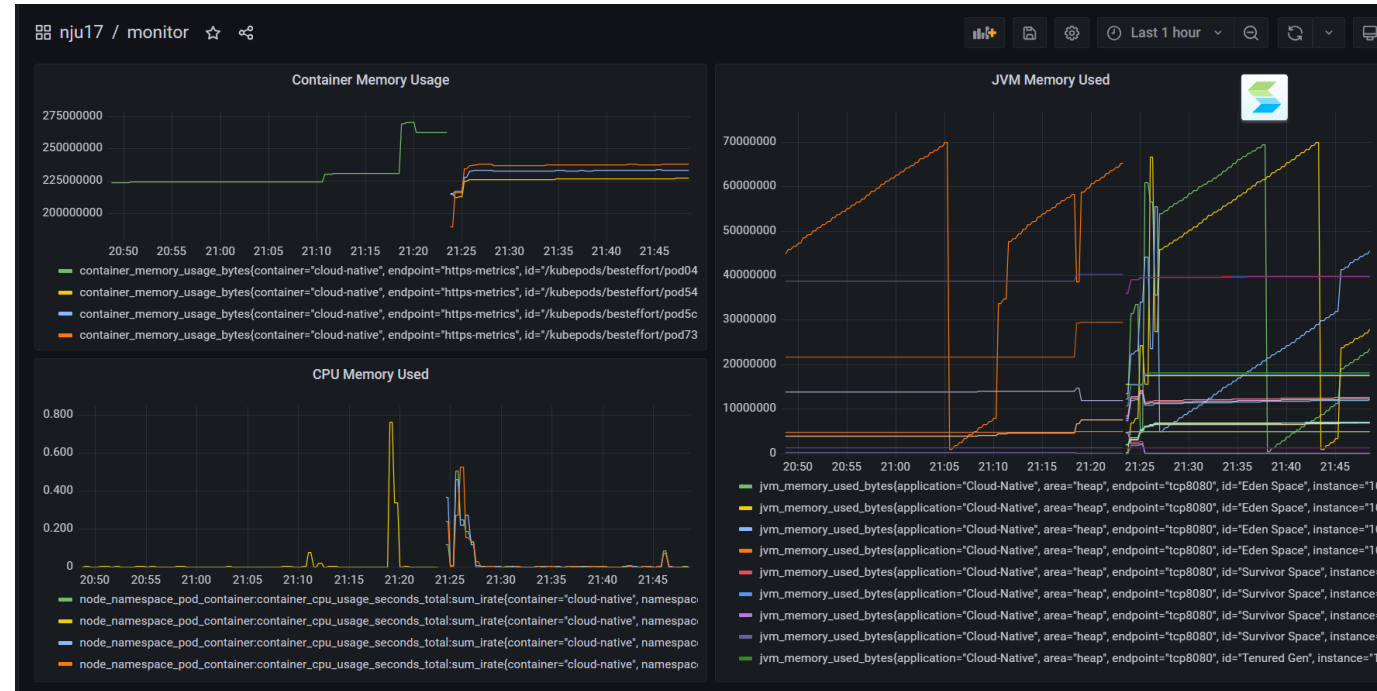
扩容后，再次对接口进行压力测试，可以看到压力测试的结果如下，请求成功的次数大大增加：



监控大屏也可以观察到相应的变化，container变为3个，曲线呈上升趋势：



扩容前后对比：



## 自动扩容

自动扩容不写了~